

NASA HW9

B11901164 陳秉緯

討論：R13941146 李毓庭, B12901194 賴睿廷

第一部分：系統環境與 NFS 基礎安裝

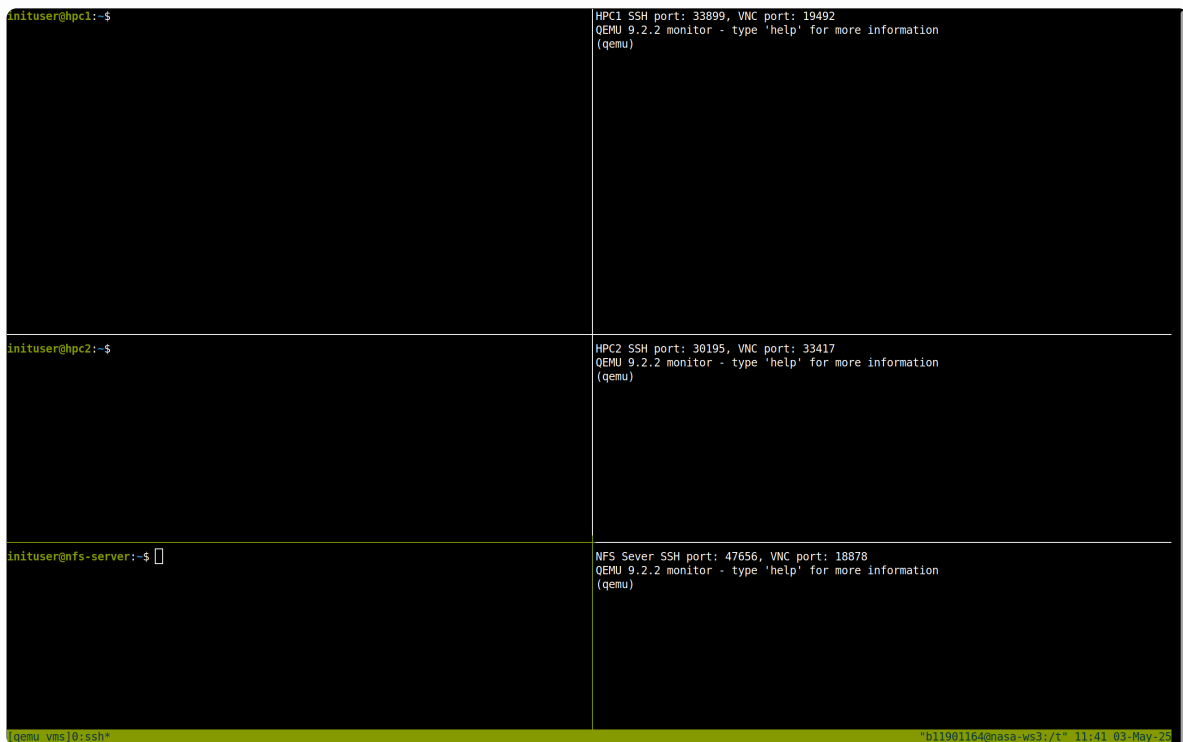
1. 環境說明

ref:

https://docs.google.com/presentation/d/12CehpBWNyFjKwv4hO1BSLZGJz_u0fkogbm1qKortKys/edit?usp=sharing

1. 啟動三台 VM：

```
cd /tmp2/b11901164  
/tmp2/hw9-release/run-vm.sh b11901164 5
```



2. 根據助教分配的 ip 分別設定 Server, hpc1, 跟 hpc2 的 ip:

- 停用 cloud-init 的網路設定：

```
sudo mkdir -p /etc/cloud/cloud.cfg.d  
echo 'network: {config: disabled}' | sudo tee /etc/cloud/cloud.cfg.d/99-  
disable-network-config.cfg
```

- 建立 netplan 設定檔：

```
sudo vim /etc/netplan/01-static-ip.yaml
```

```
network:
  version: 2
  ethernets:
    ens4:
      dhcp4: no
      addresses:
        - 192.150.9.19/16 # Server IP
      nameservers:
        addresses: [8.8.8.8, 1.1.1.1]
      routes:
        - to: default
          via: 192.150.0.1
```

- `sudo netplan apply` 套用設定
- `ip a show ens4` 確認設定是否正確

2. 題目要求

1.

On Server

1. `sudo mkdir -p /srv/nfs-share`

2. `sudo vim /etc/exports` 加入：

```
/srv/nfs-share 192.150.9.20(rw,sync,no_subtree_check)
192.150.9.21(rw,sync,no_subtree_check)
```

3. `sudo chown 1000:1000 /srv/nfs-share`
`sudo chmod 755 /srv/nfs-share`
`sudo exportfs -arv`

4. `sudo systemctl enable --now nfs-server` 啟動 nfs-server

2.

On hpc1

1. `sudo mkdir -p /mnt/nfs-share`
2. `sudo mount -t nfs 192.150.9.19:/srv/nfs-share /mnt/nfs-share`
3. `vim /mnt/nfs-share/from_hpc1.txt`：

```
Hello from hpc1
```

On hpc2

重複步驟 2-1 ~ 2-2

5. `cd /mnt/nfs-share && ls -al && cat from_hpc1.txt`

```
inituser@hpc2:~$ cd /mnt/nfs-share && ls -al && cat from_hpc1.txt
total 12
drwxr-xr-x 2 inituser inituser 4096 Apr 29 02:55 .
drwxr-xr-x 3 root      root      4096 Apr 29 02:57 ..
-rw-rw-r-- 1 inituser inituser   16 Apr 29 02:55 from_hpc1.txt
Hello from hpc1
```

6. `vim from_hpc2.txt :`

```
Hello from hpc2
```

On hpc1

7. `cd /mnt/nfs-share && ls -al && cat from_hpc2.txt`

```
inituser@hpc1:~$ cd /mnt/nfs-share && ls -al && cat from_hpc2.txt
total 16
drwxr-xr-x 2 inituser inituser 4096 Apr 29 03:11 .
drwxr-xr-x 3 root      root      4096 Apr 29 02:51 ..
-rw-rw-r-- 1 inituser inituser   16 Apr 29 02:55 from_hpc1.txt
-rw-rw-r-- 1 inituser inituser   16 Apr 29 03:11 from_hpc2.txt
Hello from hpc2
```

3. 繳交內容

- 見 2. 題目要求
- 截圖：

1. Screenshot 1:

```
inituser@nfs-server:~$ systemctl status nfs-kernel-server
● nfs-server.service - NFS server and services
   Loaded: loaded (/usr/lib/systemd/system/nfs-server.service; enabled; preset: enabled)
   Drop-In: /run/systemd/generator/nfs-server.service.d
            └─order-with-mounts.conf
   Active: active (exited) since Tue 2025-04-29 02:21:53 UTC; 54min ago
   Main PID: 876 (code=exited, status=0/SUCCESS)
     CPU: 11ms

Apr 29 02:21:52 nfs-server systemd[1]: Starting nfs-server.service - NFS server and services...
Apr 29 02:21:53 nfs-server systemd[1]: Finished nfs-server.service - NFS server and services.
```

2. Screenshot 2:

```
inituser@hpc1:~$ mount | grep nfs
192.168.167.71:/srv/nfs-share on /mnt/nfs-share type nfs4 (rw,relatime,vers=4.2,rsz=1048576,wsz=1048576,namlen=255,hard,proto=tcp,timeo=600,retrans=2,sec=sys,clientaddr=192.168.167.71,local_lock=none,addr=192.168.167.71)
inituser@hpc1:~$ df -h
Filesystem                Size      Used Avail Use% Mounted on
tmpfs                      795M    1020K    794M   1% /run
/dev/mapper/ubuntu--vg-ubuntu--lv 15G    4.4G    9.6G  32% /
tmpfs                      3.9G     0    3.9G   0% /dev/shm
shm                        5.0M     0    5.0M   0% /run/lock
/dev/sda2                  2.0G    95M    1.7G   6% /boot
tmpfs                      795M    12K    795M   1% /run/user/1000
192.168.167.71:/srv/nfs-share 15G    4.4G    9.6G  32% /mnt/nfs-share

inituser@hpc2:~$ mount | grep nfs
192.168.167.71:/srv/nfs-share on /mnt/nfs-share type nfs4 (rw,relatime,vers=4.2,rsz=1048576,wsz=1048576,namlen=255,hard,proto=tcp,timeo=600,retrans=2,sec=sys,clientaddr=192.168.167.71,local_lock=none,addr=192.168.167.71)
inituser@hpc2:~$ df -h
Filesystem                Size      Used Avail Use% Mounted on
tmpfs                      795M    1016K    794M   1% /run
/dev/mapper/ubuntu--vg-ubuntu--lv 15G    4.4G    9.6G  32% /
tmpfs                      3.9G     0    3.9G   0% /dev/shm
shm                        5.0M     0    5.0M   0% /run/lock
/dev/sda2                  2.0G    95M    1.7G   6% /boot
tmpfs                      795M    12K    795M   1% /run/user/1000
192.168.167.71:/srv/nfs-share 15G    4.4G    9.6G  32% /mnt/nfs-share
```

3. Screenshot 3:

```
inituser@nfs-server:/srv/nfs-share$ ls -al
total 16
drwxr-xr-x 2 inituser inituser 4096 Apr 29 03:11 .
drwxr-xr-x 3 root      root    4096 Apr 29 02:26 ..
-rw-rw-r-- 1 inituser inituser 16 Apr 29 02:55 from_hpc1.txt
-rw-rw-r-- 1 inituser inituser 16 Apr 29 03:11 from_hpc2.txt
inituser@nfs-server:/srv/nfs-share$

inituser@hpc2:/mnt/nfs-share$ ls -al
total 16
drwxr-xr-x 2 inituser inituser 4096 Apr 29 03:11 .
drwxr-xr-x 3 root      root    4096 Apr 29 02:57 ..
-rw-rw-r-- 1 inituser inituser 16 Apr 29 02:55 from_hpc1.txt
-rw-rw-r-- 1 inituser inituser 16 Apr 29 03:11 from_hpc2.txt
inituser@hpc2:/mnt/nfs-share$

inituser@hpc1:/mnt/nfs-share$ ls -al
total 16
drwxr-xr-x 2 inituser inituser 4096 Apr 29 03:11 .
drwxr-xr-x 3 root      root    4096 Apr 29 02:51 ..
-rw-rw-r-- 1 inituser inituser 16 Apr 29 02:55 from_hpc1.txt
-rw-rw-r-- 1 inituser inituser 16 Apr 29 03:11 from_hpc2.txt
inituser@hpc1:/mnt/nfs-share$
```

第二部分：多使用者帳號與權限控管

1. 題目要求

ref: <https://docs.google.com/presentation/d/1QOBSuBnh2F55daXRpcfpHbN-fNiUS3Hz2edsyFqzFQQ/edit?usp=sharing>, <http://old.linux.vbird.org/somepaper/20050817-lldap-1.pdf>

1.

On Server

1. `sudo apt install -y slapd ldap-utils`
2. `vim suffix.ldif`:

```
dn: olcDatabase={1}mdb,cn=config
changetype: modify
```

```
replace: olcSuffix
olcSuffix: dc=nasa,dc=csie,dc=ntu
```

3. `sudo ldapmodify -Y EXTERNAL -H ldapi:/// -f suffix.ldif` apply change

4. `vim rootdn.ldif`:

```
dn: olcDatabase={1}mdb,cn=config
changetype: modify
replace: olcRootDN
olcRootDN: cn=admin,dc=nasa,dc=csie,dc=ntu
```

5. `sudo ldapmodify -Y EXTERNAL -H ldapi:/// -f rootdn.ldif` apply change

6. `slappasswd` and copy the result

7. `vim rootpw.ldif`:

```
dn: olcDatabase={1}mdb,cn=config
changetype: modify
replace: olcRootPW
olcRootPW: {SSHA}xiOeZFSc5hUpkNaQn6n0MEiYhYBfLjXe
```

8. `sudo ldapmodify -Y EXTERNAL -H ldapi:/// -f rootpw.ldif` apply change

9. `vim base.ldif`:

```
dn: dc=nasa,dc=csie,dc=ntu
dc: nasa
objectClass: top
objectClass: domain

dn: cn=admin,dc=nasa,dc=csie,dc=ntu
cn: admin
objectClass: organizationalRole
description: admin account

dn: ou=people,dc=nasa,dc=csie,dc=ntu
ou: people
objectClass: organizationalUnit

dn: ou=group,dc=nasa,dc=csie,dc=ntu
ou: group
objectClass: organizationalUnit
```

10. `sudo ldapadd -D cn=admin,dc=nasa,dc=csie,dc=ntu -W -H ldapi:/// -f base.ldif` 輸入 password

但是出現錯誤

ldap_bind: Invalid credentials (49)

11. `sudo dpkg-reconfigure slapd` 重新設定password: nasa2025 然後再做一次步驟 10，就可以了

12. `vim group.ldif`:

```
dn: cn=astro,ou=group,dc=nasa,dc=csie,dc=ntu
objectClass: posixGroup
cn: astro
gidNumber: 20000
```

13. `sudo ldapadd -x -D cn=admin,dc=nasa,dc=csie,dc=ntu -W -f group.ldif` **apply**
change

14. add astro1, astro2, and astro3:

- `vim astro1.ldif:`

```
dn: uid=astro1,ou=people,dc=nasa,dc=csie,dc=ntu
objectClass: inetOrgPerson
objectClass: posixAccount
objectClass: shadowAccount
cn: astro1
sn: astro1
uid: astro1
uidNumber: 20001
gidNumber: 20000
homeDirectory: /home/astro1
loginShell: /bin/bash
userPassword: {SSHA}xiOeZFSc5hUpkNaQn6n0MEiYhYBfLjXe
```

- `sudo ldapadd -x -D cn=admin,dc=nasa,dc=csie,dc=ntu -W -f astro1.ldif` **apply**
change

- `vim astro2.ldif:`

```
dn: uid=astro1,ou=people,dc=nasa,dc=csie,dc=ntu
objectClass: inetOrgPerson
objectClass: posixAccount
objectClass: shadowAccount
cn: astro2
sn: astro2
uid: astro2
uidNumber: 20002
gidNumber: 20000
homeDirectory: /home/astro2
loginShell: /bin/bash
userPassword: {SSHA}xiOeZFSc5hUpkNaQn6n0MEiYhYBfLjXe
```

- `sudo ldapadd -x -D cn=admin,dc=nasa,dc=csie,dc=ntu -W -f astro2.ldif` **apply**
change

- `vim astro3.ldif:`

```
dn: uid=astro1,ou=people,dc=nasa,dc=csie,dc=ntu
objectClass: inetOrgPerson
objectClass: posixAccount
objectClass: shadowAccount
cn: astro3
sn: astro3
uid: astro3
```

```
uidNumber: 20003
gidNumber: 20000
homeDirectory: /home/astro3
loginShell: /bin/bash
userPassword: {SSHA}xiOeZFSc5hUpkNaQn6n0MEiYhYBfLjXe
```

- `sudo ldapadd -x -D "cn=admin,dc=nasa,dc=csie,dc=ntu" -W -f astro3.ldif`
apply change

On hpc1 and hpc2

15. `sudo apt install libnss-ldapd libpam-ldapd nslcd`

16. `sudo vim /etc/nslcd.conf` 改成以下內容：

```
uri ldap://192.150.9.19/ base dc=nasa,dc=csie,dc=ntu binddn
cn=admin,dc=nasa,dc=csie,dc=ntu bindpw nasa2025
```

17. `sudo vim /etc/nsswitch.conf` 改成以下內容：

```
passwd: files systemd ldap group: files systemd ldap shadow: files systemd ldap
```

18. `sudo pam-auth-update` 選建立家目錄後按 Enter，但是後來發現登入 astro1 之後 create /home/astro1 失敗，所以 `sudo vim /etc/pam.d/common-session` 加入：

```
session required pam_mkhomedir.so skel=/etc/skel umask=0022
```

19. `sudo systemctl restart nslcd` restart

```
inituser@hpc1:~$ su - astro1
Password:
Creating directory '/home/astro1'.
astro1@hpc1:~$ su - astro2
20. Password:
Creating directory '/home/astro2'.
astro2@hpc1:~$ su - astro3
Password:
Creating directory '/home/astro3'.
```

```
astro1@hpc1:~$ su - astro1
Password:
Creating directory '/home/astro1'.
astro1@hpc2:~$ su - astro2
Password:
Creating directory '/home/astro2'.
astro2@hpc2:~$ su - astro3
Password:
Creating directory '/home/astro3'.
```

On Server,

21. `cd /srv/nfs-share`

22. `sudo mkdir astro1_dir astro2_dir astro3_dir` create them

23. 設定權限：

```
sudo chown 20001:20000 astro1_dir sudo chown 20002:20000 astro2_dir sudo chown
20003:20000 astro3_dir sudo chmod 700 astro1_dir sudo chmod 700 astro2_dir sudo
chmod 700 astro3_dir
```

24. `sudo vim /etc/exports` 改成以下：

```
! /srv/nfs-share 192.150.9.20(rw,sync,no_subtree_check,root_squash)
192.150.9.21(rw,sync,no_subtree_check,root_squash)
```

25. 更新：

```
sudo exportfs -arv sudo systemctl restart nfs-server
```


2. 繳交內容

- 見 1. 題目要求
- 截圖：

1. Screenshot 1:

```
inituser@hpc1:~$ id astro1
uid=20001(astro1) gid=20000(astro) groups=20000(astro)
inituser@hpc1:~$ id astro2
uid=20002(astro2) gid=20000(astro) groups=20000(astro)
inituser@hpc1:~$ id astro3
uid=20003(astro3) gid=20000(astro) groups=20000(astro)
```

2. Screenshot 2:

```
inituser@nfs-server:/srv/nfs-share$ ls -al
total 28
drwxr-xr-x 5 inituser inituser 4096 May  3 08:32 .
drwxr-xr-x 3 root      root      4096 May  3 05:57 ..
drwx----- 2      20001      20000 4096 May  3 08:32 astro1_dir
drwx----- 2      20002      20000 4096 May  3 08:32 astro2_dir
drwx----- 2      20003      20000 4096 May  3 08:32 astro3_dir
-rw-rw-r-- 1 inituser inituser   16 May  3 06:01 from_hpc1.txt
-rw-rw-r-- 1 inituser inituser   16 May  3 06:28 from_hpc2.txt
```

3. Screenshot 3:

```
astro1@hpc1:/mnt/nfs-share/astro1_dir$ vim from_hpc1.txt
astro1@hpc1:/mnt/nfs-share/astro1_dir$ ls
from_hpc1.txt
astro1@hpc1:/mnt/nfs-share/astro1_dir$ cat from_hpc1.txt
Hello from hpc1
```

```
inituser@hpc1:/mnt/nfs-share$ ls
astro1_dir astro2_dir astro3_dir from_hpc1.txt from_hpc2.txt
inituser@hpc1:/mnt/nfs-share$ cd astro1_dir
-bash: cd: astro1_dir: Permission denied
inituser@hpc1:/mnt/nfs-share$ su - astro2
Password:
astro2@hpc1:~$ cd /mnt/nfs-share/
astro2@hpc1:/mnt/nfs-share$ ls
astro1_dir astro2_dir astro3_dir from_hpc1.txt from_hpc2.txt
astro2@hpc1:/mnt/nfs-share$ cd astro1_dir
-bash: cd: astro1_dir: Permission denied
astro2@hpc1:/mnt/nfs-share$ su - astro3
Password:
astro3@hpc1:~$ cd /mnt/nfs-share/
astro3@hpc1:/mnt/nfs-share$ ls
astro1_dir astro2_dir astro3_dir from_hpc1.txt from_hpc2.txt
astro3@hpc1:/mnt/nfs-share$ cd astro1_dir
-bash: cd: astro1_dir: Permission denied
astro3@hpc1:/mnt/nfs-share$
```

第三部分：效能與大規模檔案測試

ref: <https://blog.gtwang.org/linux/dd-command-examples/>,
<https://stackoverflow.com/questions/28353409/bash-format-uptime-to-show-days-hours-minutes>

1. NFS 參數：

On Server

1. `sudo vim /etc/exports`：

◦ for A and B:

```
/srv/nfs-share 192.150.9.20(rw,sync,no_subtree_check,root_squash) 192.150.9.19
```

◦ for C and D:

```
/srv/nfs-share 192.150.9.20(rw,async,no_subtree_check,root_squash) 192.150.9.19
```

2. restart:

```
sudo exportfs -arv
sudo systemctl restart nfs-server
```

On hpc1 and hpc2

3. `sudo umount /mnt/nfs-share`

4. for A and C: `sudo mount -t nfs -o rsize=8192,wsiz=8192 192.150.9.19:/srv/nfs-share /mnt/nfs-share`

for B and D: `sudo mount -t nfs -o rsize=32768,wsiz=32768 192.150.9.19:/srv/nfs-share /mnt/nfs-share`

2. Shell Script：

5. `sudo vim nfs_test_b11901164.sh`：

```
#!/bin/bash

# env
MOUNT_DIR="/mnt/nfs-share"
FILENAME="${USER}_dir/testfile_${USER}_$(hostname)"
FILEPATH="$MOUNT_DIR/$FILENAME"

# remove the existing file
if [ -f "$FILEPATH" ]; then
    rm -f "$FILEPATH"
fi

# get the load before testing
```

```

LOAD_BEFORE=$(uptime | awk -F'load average: ' '{print $2}' | cut -d',' -f1)

echo "建立測試檔案：$FILEPATH"

# write 1GB test file
START_WRITE=$(date +%s.%N)
dd if=/dev/zero of=$FILEPATH bs=1M count=1024 oflag=direct status=none
END_WRITE=$(date +%s.%N)
WRITE_TIME=$(echo "$END_WRITE - $START_WRITE" | bc)
WRITE_SPEED=$(echo "scale=2; 1024 / $WRITE_TIME" | bc)

# Read
START_READ=$(date +%s.%N)
dd if=$FILEPATH of=/dev/null bs=1M iflag=direct status=none
END_READ=$(date +%s.%N)
READ_TIME=$(echo "$END_READ - $START_READ" | bc)
READ_SPEED=$(echo "scale=2; 1024 / $READ_TIME" | bc)

# avg speed
AVG_SPEED=$(echo "scale=2; ($WRITE_SPEED + $READ_SPEED) / 2" | bc)

# get the load after testing
LOAD_AFTER=$(uptime | awk -F'load average: ' '{print $2}' | cut -d',' -f1)

# output
echo "寫入時間：$WRITE_TIME s"
echo "寫入速率：$WRITE_SPEED MB/s"
echo "讀取時間：$READ_TIME s"
echo "讀取速率：$READ_SPEED MB/s"
echo "平均傳輸速率：$AVG_SPEED MB/s"
echo "CPU 使用率/負載：$LOAD_BEFORE → $LOAD_AFTER"

```

6. 如果是單一使用者則只要在hpc1執行 `nfs_test_b11901164.sh`，如果是多使用者則分別在 hpc1 與 hpc2 內的 astro1 ~ astro3 執行 `nfs_test_b11901164.sh`

```

sudo chmod +x nfs_test_b11901164.sh
./nfs_test_b11901164.sh

```

7. shell script 的運作流程以及所使用的工具：一開始先刪除舊的測試檔，再利用 dd 工具寫入 1GB 檔案至 /mnt/nfs-share，然后再讀取該檔案，分別計算寫入與讀取的時間與速率。透過 uptime 擷取測試前後的 CPU 負載。

8. 表格：

測試情境	寫入時間 (s)	讀取時間 (s)	平均傳輸速率 (MB/s)	CPU 使用率/ 負載
組合 A + 單一使用者	189.82	13.13	41.67	1.88
組合 A + 多使用者	850.58	54.61	10.44	4.24
組合 B +	205.44	10.38	51.80	1.36

單一使用者				
組合 B + 多使用者	614.23	55.19	10.11	3.54
組合 C + 單一使用者	27.21	22.54	41.52	1.72
組合 C + 多使用者	139.40	66.22	11.18	4.62
組合 D + 單一使用者	15.58	17.85	61.53	0.62
組合 D + 多使用者	108.65	58.71	13.40	3.71

9. 問答題：

- 有，在多使用者同時讀寫的情況下效能明顯下降。以組合 A 為例，單一使用者平均傳輸速率為 41.67 MB/s，而多使用者情境下降至 10.44 MB/s，其他組合也呈現類似趨勢。
- 原因：
 - 多使用者同時存取導致網路擁塞與 NFS 鎖定 (lock contention)
 - sync 模式會強制等待資料寫入磁碟，造成延遲更加明顯
 - rsize/wsize 值越小，封包數越多，傳輸效率更低
 - CPU 負載上升也影響 I/O 效能
- 解決辦法：使用 async 並適當增加 rsize/wsize（如組合 D），能顯著改善效能；若資料一致性不是關鍵，可優先考慮 async 模式，並使用更快的網路介面，如 10GbE 以減少瓶頸。