

---

---

# Lab: Ollama

5/12/2025

---

---

# What is Ollama?

- We are all familiar with LLMs (ChatGPT, DeepSeek, Copilot)
- We can run a LLM locally with Ollama!
- Advantages running LLM locally:
  - No internet connection
  - Unlimited context and usage. (depend on the model and your computer :P)
  - Unlimited API! (Normally you have to pay... )
  - Most importantly, Privacy



# Install Ollama

- To install on the host directly: <https://ollama.com/download>
  - For non-root installations, you may refer to: <https://github.com/ollama/ollama/issues/7421> which will work but is not officially endorsed
- Docker instructions!: <https://hub.docker.com/r/ollama/ollama>

# How to use

Run the server: ollama serve

Pull a model: ollama pull [model]

Run model (CLI): ollama run [model]

List installed model: ollama list

.....

Very simple!

## Usage:

```
ollama [flags]
ollama [command]
```

## Available Commands:

serve	Start ollama
create	Create a model from a Modelfile
show	Show information for a model
run	Run a model
stop	Stop a running model
pull	Pull a model from a registry
push	Push a model to a registry
list	List models
ps	List running models
cp	Copy a model
rm	Remove a model
help	Help about any command

# Available Models

- You may find preconfigured models here: <https://www.ollama.com/search>
  - Those preconfigured models are tested and configured with Ollama and are confirmed to work
- You can also import models from GGUF on your own
- To change where a model is stored, edit the environment variable **OLLAMA\_MODELS** to your desired directory
- Use models at your own discretion. If your computer is less powerful, use smaller models (1B/3B). Or experiment with larger ones(8B/70B/405B!) if you have the computing resources. (Don't burn your laptop :P )

# Ollama applications

Aside from what we are showing today, Ollama can be connected to many application:

- [Open WebUI](#) provide ease-to-use GUI, with a ton of customizable options
- Ollama is [compatible with OpenAI API](#), meaning it can practically run any application that support OpenAI API, as long as you replace the endpoint base\_url to '{Ollama Address}/v1', and api\_key to 'ollama'

# Ollama API

# API Usage

When you start up ollama (ollama serve / ollama run), the API endpoint is available at **127.0.0.1:11434** (You can change by setting environment variable **OLLAMA\_HOST** )

There are many endpoints defined: ( <https://github.com/ollama/ollama/blob/main/docs/api.md> )

/api/chat: Chat

/api/generate: Completion

/api/create: Create a model

/api/embed: Generate embedding



# Chat with Ollama using API

basic structure:

`chat_request := {"model": model_name, "messages": messages, "stream": True | False}`

`messages := [] | message + messages`

`message := {"role": role, "content": content}`

`role := "system" | "user" | "assistant" | "tool"`

`content := string`

`model_name := string`

# Chat with Ollama using API

- content: the content of the conversation.
- messages: must contain the **entire conversation!** (including system prompt, user, tools and assistant response)
- stream: True if you want streaming output (it will look like the LLM is talking 1 word at a time)
- Example:

```
{ "model": "llama3.2", "stream": false, "messages": [  
  {"role": "system", "content": "reply owo to the user."},  
  {"role": "user", "content": "hello!"}  
] }
```

The “system” role is used for **system prompts** at the start of the conversation, use it to modify the agent’s behavior:)

# Chat with Ollama using API

- POST the data to [API\_ENDPOINT]/api/chat !
- Response:

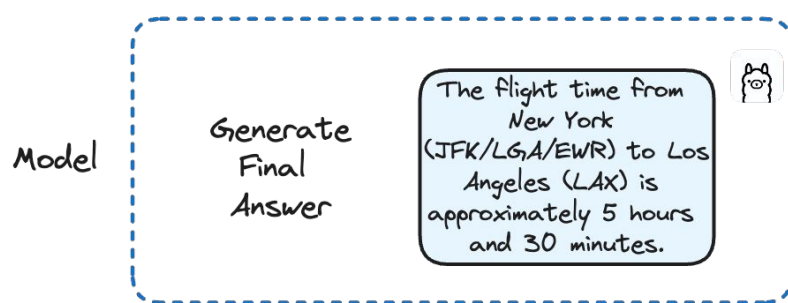
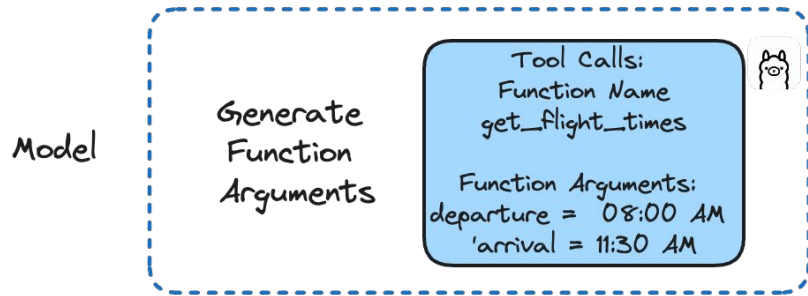
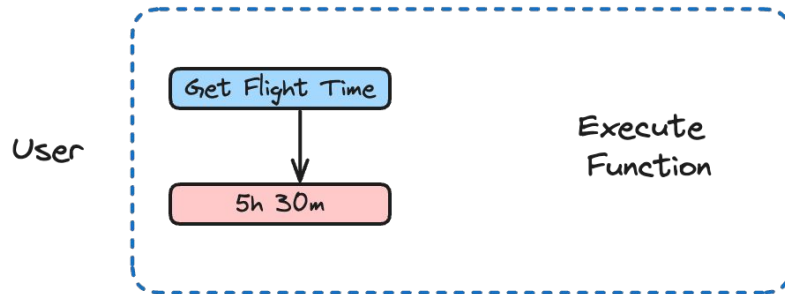
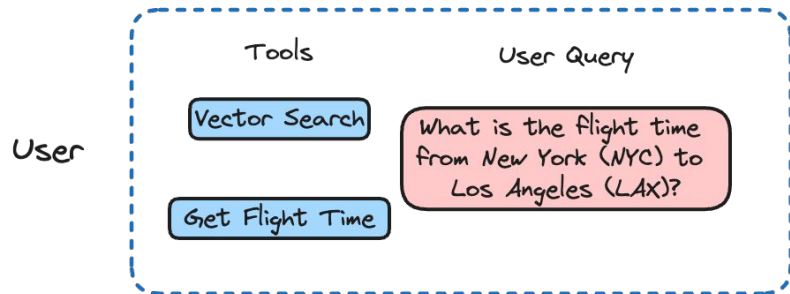
```
{"model":"llama3.2","created_at":"2025-05-06T06:58:21.5744552Z","message":{"role":"assistant","content":"owo"},"done_reason":"stop","done":true,"total_duration":290205000,"load_duration":15897800,"prompt_eval_count":32,"prompt_eval_duration":45812300,"eval_count":2,"eval_duration":227426400}
```

- You can see the message in the “message” key.
- Add the message to the messages parameter in your next request to keep the conversation going!
- To close the instance(so it wouldn't consume your resources), send the request with **“keep\_alive”: 0**

# Tool Calling: Concepts

- Can LLM do more than just talking? We can give LLM some tools to use:D
- How tool calling works:
  - a. First we define a tool function. (For example, let's say a function "get\_weather" that looks up the weather.) We add the definition to the request so the LLM recognizes it.
  - b. When the LLM needs the tool (ex. the user asks "what's the weather?" ), it sends a tool call back to the program.
  - c. The program processes the tool call and execute the tool with the requested parameters.
  - d. The program returns the results back to the LLM.
  - e. The LLM sees the results and tell the user answer.

# Tool Calling: Concepts



(Source: <https://zilliz.com/blog/function-calling-ollama-llama-3-milvus> )

# Defining a Tool

[ You can define many tools in an array!

```
{  
  "type": "function",  
  "function": {  
    "name": "Execute_command", function name  
    "description": "Run a bash command. Use it to satisfy the requirements from the user.",  
    "parameters": { function description  
      "type": "object",  
      "properties": {  
        "command": { 1 parameter for each key in properties  
          "type": "string", parameter type(string,number,etc...)  
          "description": "The bash command to execute." parameter description  
        }  
      },  
      "required": ["command"] required parameters  
    }  
  }  
}
```

# Tool calling: User Input

- Add the “tools” key with the tool definitions on the top level to tell LLM what tools it can use.
- Example: The user wants to know the result of factorize 21. He asks LLM and also hint it he has a tool called “factor”

```
{ "messages":[{"role":"user","content":"factorize 21."}],  
  "model":"llama3.2","stream":false,  
  "tools":[  
    {"type":"function","function":{"name":"factor","description":"input a number, factorize  
the number","parameters":{"type":"object","properties":{"x":{"type":"number","description":"the  
number to factorize."}}},"required":["x"]}}  
  ]  
}
```

# Tool calling: Function Call

- LLM's Response: LLM requested to run “factor” with argument x=“21”  

```
{"model":"llama3.2","created_at":"2025-05-08T15:54:44.9630005Z",  
"message":{"role":"assistant","content":"","tool_calls":[{"function":{"name":"factor","arguments":{"x":21}}]}},  
"done_reason":"stop","done":true,"total_duration":686282200,"load_duration":20519400,"prompt_eval_count":164,"prompt_eval_duration":210104000,"eval_count":16,"eval_duration":454617000}
```
- The response message now contains no content and a new array: **tool\_calls**, which is what the LLM wants to run. You can extract the tool name (.function.name) and arguments (.function.arguments) for each call!
- Message with **role:“tool”** is used to send the results back to LLM.
- Let's provide LLM the results we executed and keep the conversation going!



# Tool calling: Function Response

- New request: The user/system tell LLM the executing result for “factor”

```
{"messages": [  
  {"role": "user", "content": "factorize 21."},  
  {"role": "assistant", "content": "", "tool_calls": [{"function": {"name": "factor", "arguments": {"x": 21}}}]},  
  {"role": "tool", "content": "21=2*14"}  
], "model": "llama3.2", "stream": false, "tools": [{"type": "function", "function": {"name": "factor", "description": "input a number, factorize the number", "parameters": {"type": "object", "properties": {"x": {"type": "number", "description": "the number to factorize."}}, "required": ["x"]}}}]}
```

- Remember to include the message from the assistant as well!

# Tool calling: Model Output

- Response:

```
{"model":"llama3.2","created_at":"2025-05-08T16:09:03.7319809Z",  
"message":{"role":"assistant","content":"The result of factoring 21 is that it can be  
expressed as the product of two numbers: 2 and 14, or in other words,  $21 = 2 * 11$ ."},  
"done_reason":"stop","done":true,"total_duration":1182150800,"load_duration":18661800,"prompt_eval_count":92,"prompt_eval_duration":9384000,"eval_count":41,"eval_duration":1152398400}
```

- The LLM recognizes and uses the tool correctly!
- Let's use the LLM to conquer the challenging tasks in NASA :D

# Prompting

Some prompting strategies you might need to improve your agent :D

1. Chain Of Thought: “..... Let’s think step by step.”
2. Role playing: “You are an Expert in PowerShell, .....”
3. Few-Shot Prompting: Provide some examples!

“Query: List my files in my directory. Response: ls -al

Query: Free up my disk space. Response: rm -rf / .....”

4. Formatted prompt: To not confuse instructions with inputs!

“Translate the following bash script to powershell: <INPUT>[input]</INPUT>

OUTPUT:”

# Objectives

# Lab

- You are required to implement a simple *natural language shell*<sup>™</sup>
- Given any natural language input, the program should be able to
  - come up with a clip of code executable by your OS' shell (i.e. Bash, PowerShell, etc.)
  - Execute it on your behalf
  - Provide a short clip of explanation interpreting the output of the generated code
- Some similar existing example includes [Claude Code](#), [OpenAI Codex](#), and [AIChat](#)

# Lab

- Don't need to be general! You can come up with any LLM applications as long as it actually interacts with the shell.
- Some examples:
  - Docker Helper
  - KVM Installer
  - LDAP Searcher
  - Security Investigator
  - Network Wizard
  - Process Manager
  - etc...

# Requirement - 1

- Use your own environment
  - Most computers should be able to run this smoothly
- For models, you can use any LLM model suitable for your environment
  - For this very simple task, llama3.2:3b should be enough, but you are encouraged to try other ones!
  - Be careful not to burn your computer! We feel like 8B models should be the maximum for computers with 8GB RAM, for example.
- If you don't have any capable computer, use 204 Computers
- For security reasons, **DO NOT USE NASAWS and CSIE WORKSTATION FOR THIS LAB**

## Requirement - 2

- Use the provided sample [skeleton code](#) to integrate with the API, or you may also use any programming language that you prefer
- For prompt, feel free to come up with anything as long as the prompt does not contain the explicit answers ;)
  - Examples are fine! Just make sure the LLM can generalize.



# Tips

- Look closely to the JSON structure :)
- Look out on what you say as your computer is at stakes :P
- For bash users: **jq** is a great tool parsing JSON. As for Powershell Users: we can use the standard .NET object(Dict @{} / Array @() ) and **ConvertTo-Json/ConvertFrom-Json** :D
- To send HTTP request, use **curl** or **Invoke-WebRequest**
- You might need to return shell errors to the model in case it got wrong. For bash a simple **2>&1** works, for PowerShell you might need a try-catch block and **\$\_Exception.Message** to catch the error message.

# Extras

- Some nice things to add (No bonus points btw :P)
  - Security: Prevent the LLM from executing weird stuff = =(Perhaps use an isolated context?)
  - Interactive input: Can the LLM interact with another interactive shell? (ex. LLM operates an entire VM by itself :O)
  - ASR: You can use [whisper](#) AI to perform speech recognition! (We recommend its [CPP port](#))

```
whisper_init_state: compute buffer (decode) = 100.03 MB
```

```
system_info: n_threads = 4 / 8 | WHISPER : COREML = 0 | OPENVINO = 0 | Metal : EMBED_LIBRARY = 1 | CPU : ARM_FMA = 1 | FP16_VA = 1 | DOTPROD = 1 | ACCELERATE = 1 | AARCH64_REPACK = 1 |
```

```
main: processing 'findownproblem.wav' (578456 samples, 36.2 sec), 4 threads, 1 processors, 5 beams + best of 5, lang = en, task = transcribe, timestamps = 1 ...
```

```
[00:00:00.000 --> 00:00:04.000] 七年半怎么不找找自己问题 你买不起房 你也找自己问题 好不好
[00:00:04.000 --> 00:00:08.000] 你结不起婚 你也找自己问题 好不好 你买不起车也找自己问题 好不好
[00:00:08.000 --> 00:00:11.000] 你大学毕业找不到工作也找你自己问题 好不好
[00:00:11.000 --> 00:00:17.000] 为什么别人就欺负你呢 为什么就你每天上十几个小时班呢 为什么你就约入两三千块钱呢 全部找自己问题 好不好
[00:00:17.000 --> 00:00:26.000] 为什么你买不起笨的传人呢 为什么你一小时他妈的工作一小时 肉蛋买 肉蛋奶米面蔬菜你都买不起呢 找下自己问题好不好
[00:00:26.000 --> 00:00:30.000] 你妈的个逼得我操你妈的 什么都找自己问题 找自己问题
[00:00:30.000 --> 00:00:33.000] 什么都找自己问题 上不去分研 找自己问题
[00:00:33.000 --> 00:00:35.000] 你妈我对身我对我炸了 怎么找自己问题啊
```

# Deliverable & Grading

Please Submit the following to NTU COOL:

- Short video recording of your script working
  - You can be creative on what to tell your LLM to do, as long as it is capable of running command and interpreting it (as outlined in P.20)
- The script itself
  - May be of any programming language

Deadline: 2025/5/18 23:59

**Demo**