

Linux Kernel & Related Topics

Michael Tsai
2025/4/21

Agenda

- Before class:
please take a look at various readings on NTU COOL
 - Software licenses
 - The U. Minnesota incident
- Lab 9: LDAP
(Light-weight Directory Access Protocol)
- HW8 (LDAP) will be announced today.

What does Kernel do?

- Kernel creates these concepts from the low-level hardware features:
 - Processes (time-sharing, process address space)
 - Signals and semaphores
 - Virtual memory (swapping, paging, mapping)
 - The filesystem (files, directories, namespace)
 - General input/output (specialty hardware, keyboard, mouse, USB, etc.)
 - Interprocess communication (pipes and network connections)

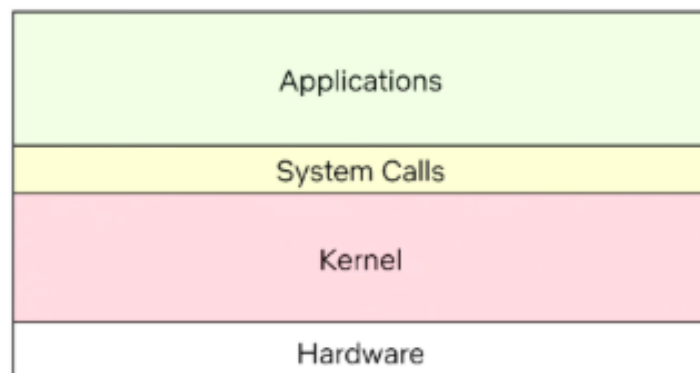
Linux kernel in the past 20 years (FOSDEM 2020)

- SMP support from v2.0
 - Big kernel lock:
lock_kernel() & unlock_kernel() - v2.4 「暫時」的鎖定機制
 - BKL用於保護整個核心、spinlock保護特定的某共享資源
- Support for virtualization
 - KVM == Kernel-based Virtual Machine (AWS一開始用Xen, 後AWS & Google Cloud皆使用KVM)
 - x86架構非為虛擬化而生，因此軟體架構複雜
 - 直接把Linux轉為hypervisor，結合硬體的虛擬化支援
- Data Plane Development Kit
 - app-to-wire-to-app, TCP/UDP的延遲到達 6 us
 - 繞過作業系統的管理，使得應用程式可以直接操作、寫入網路介面緩衝區
 - 攔截libc呼叫、確認後把呼叫轉到使用者層級的TCP/IP stack
- Ref: jserv <http://hackmd.io/@sysprog/linux-dev-review>

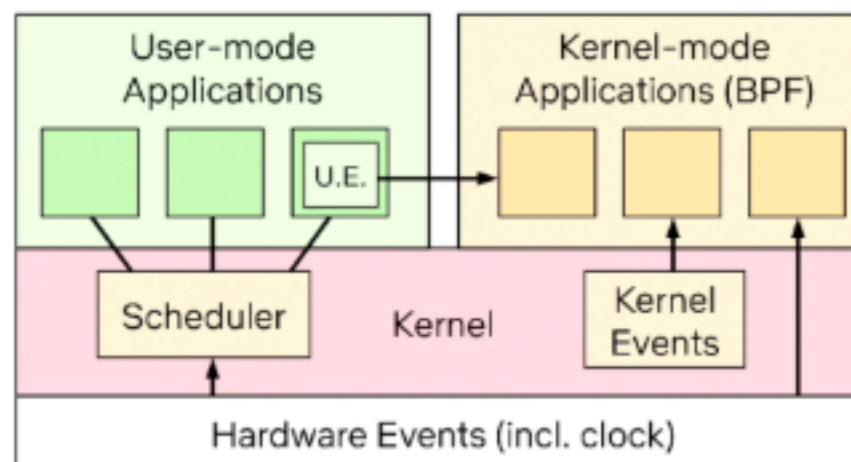
Linux kernel in the past 20 years (FOSDEM 2020)

- Asynchronous I/O:
 - Async I/O: 從/往核心空間讀取或寫入資料的過程，交由核心內部處理。複製完資料後，再通知使用者的行程。這中間使用者行程不會被阻擋。
 - 更有效率、速度更快
- Container的支援：
 - Control groups (cgroups) —> 限制，控制與隔離 Process 所使用到的系統資源 [ref]
 - Namespaces —> 命名空間 (Namespaces) 是 Linux 核心的一項功能，它能夠將核心資源進行劃分，使得一組行程 (processes) 看到的是一組資源，而另一組行程則看到的是另一組不同的資源。
- Microkernel化

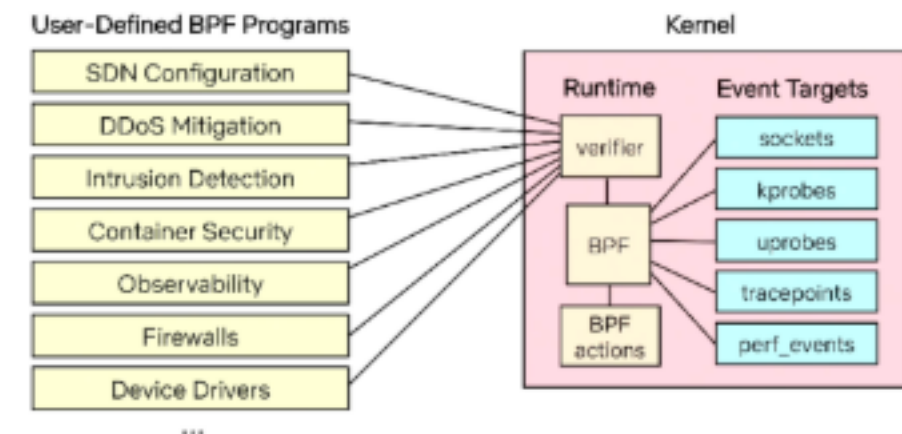
50 Years, one (dominant) OS model



Modern Linux: Event-based Applications

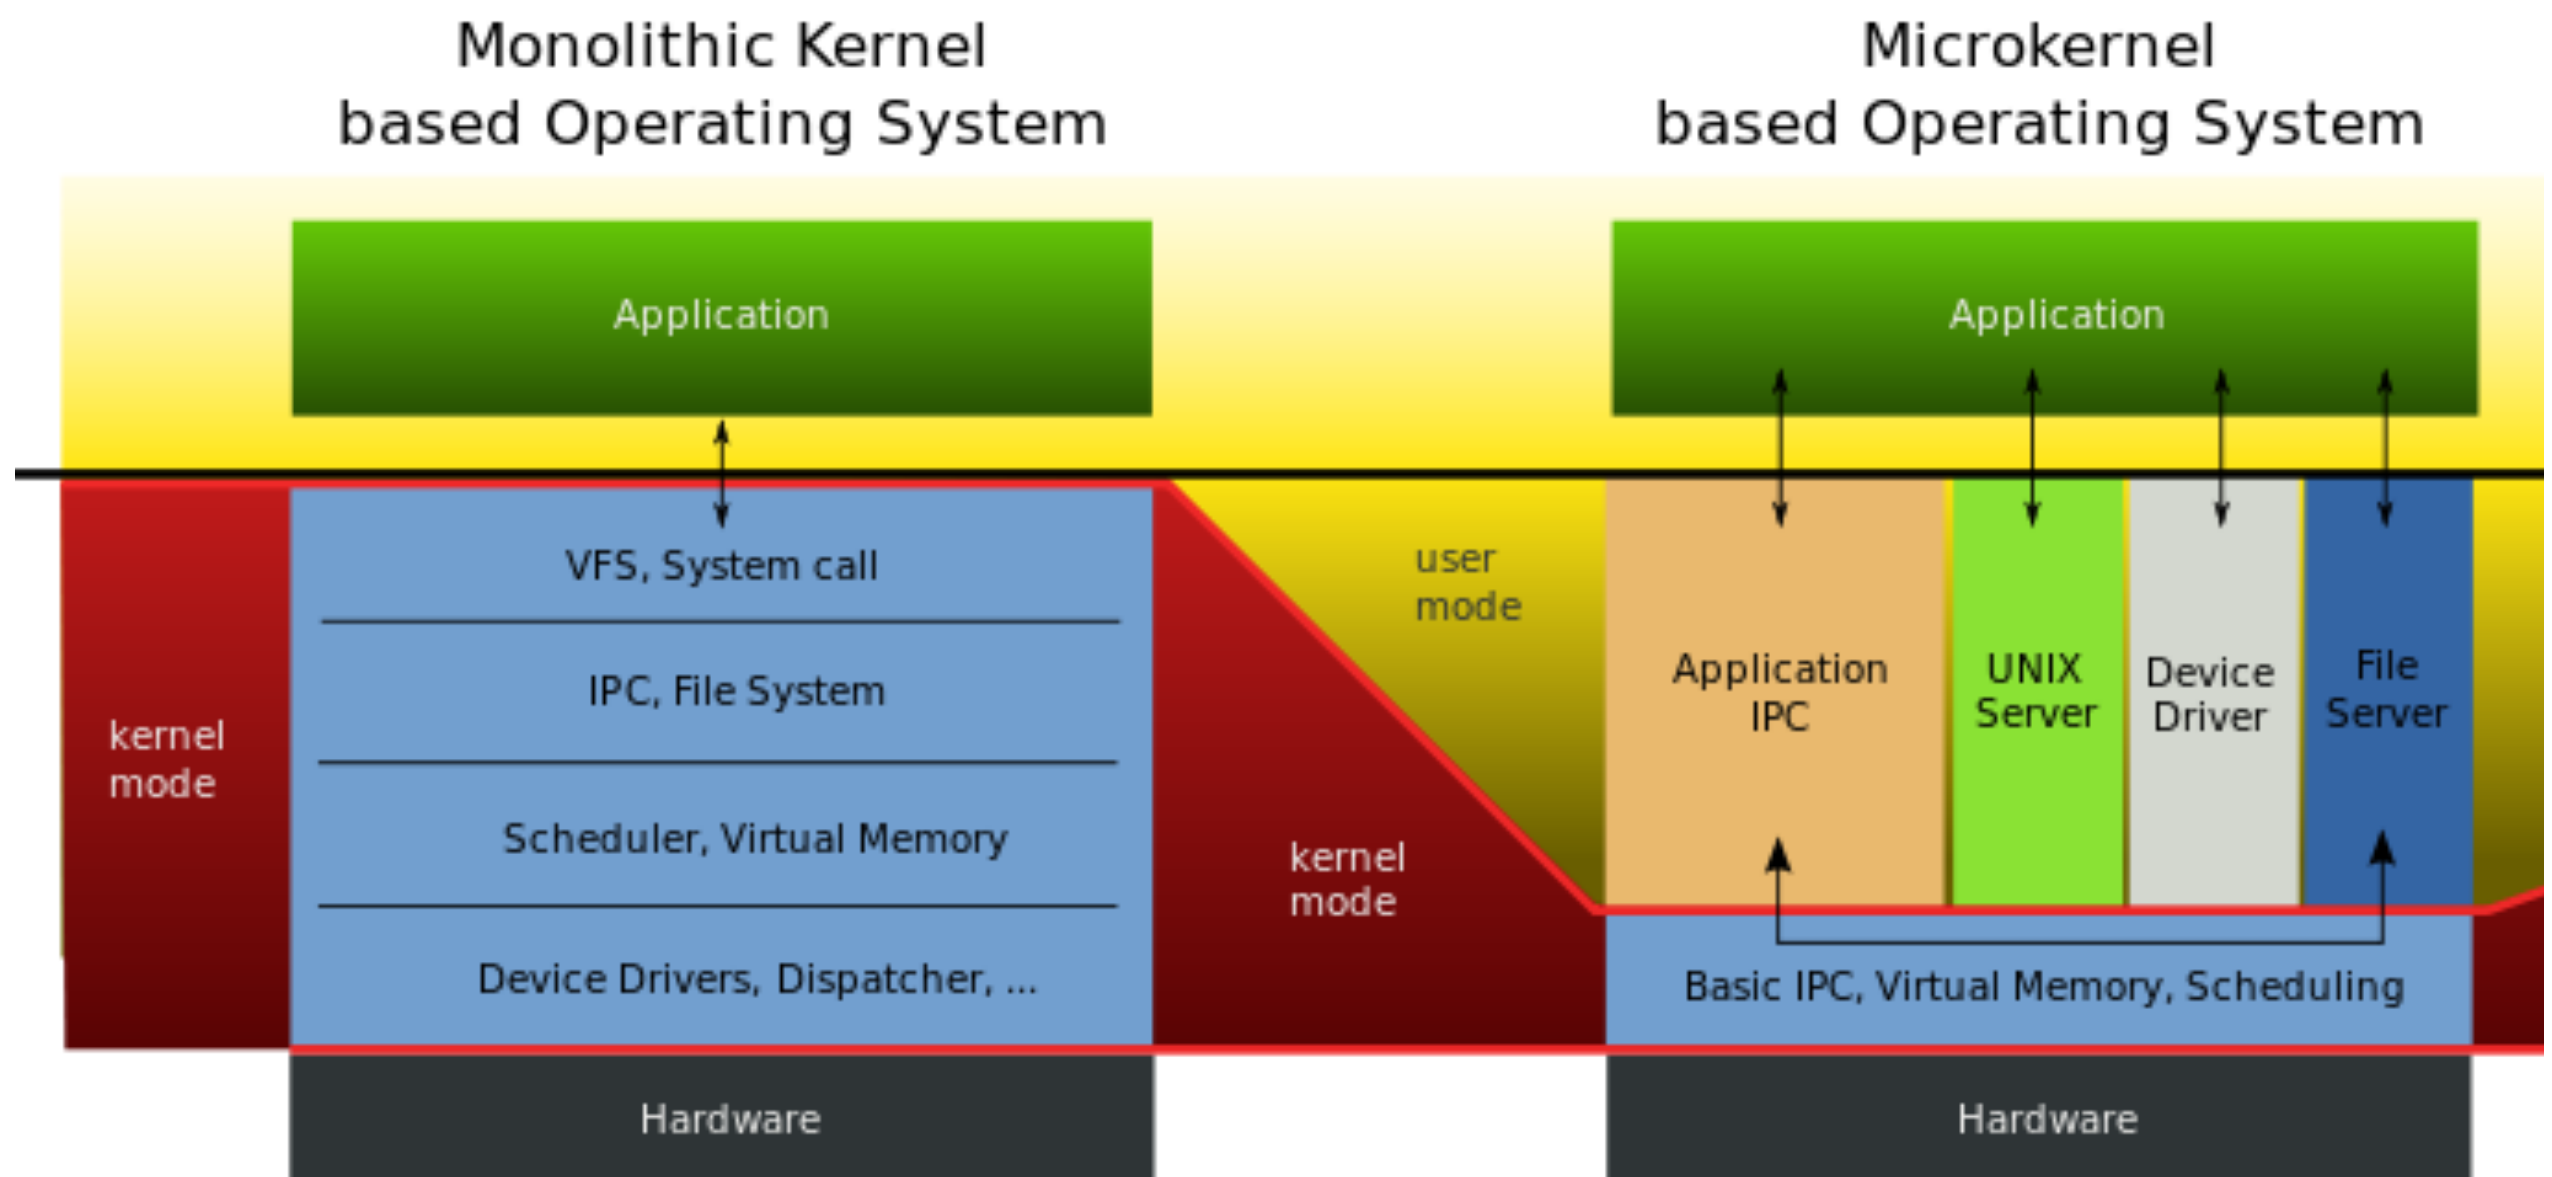


BPF 2019



Kernel Adaptation

- Linux is a monolithic kernel at heart (at the beginning)
 - Monolithic kernel: the entire OS runs in **kernel space** - a section of memory reserved for privileged operating system functions (device drivers, IPC, virtual memory, scheduling all run in the same address space)
 - Microkernel: many of the “services” run in user mode as regular processes
- Modern monolithic kernels support on-demand loading of modules
 - No need to re-build the entire kernel & **reboot**
 - Example: filesystem and device drivers
- Modern OS'es (e.g., Windows, Mac OS X) use hybrid kernel
- Linux is moving toward having microkernel



GNU

General Public License (GPL)

- Find online about GPL, L-GPL & BSD licenses. What are their major differences?
- Since version 0.12, Linux kernel is released under GPL. How has the use of GPL helped the development of Linux?
- Choose a license
<https://choosealicense.com/licenses/>
- Read:
GPL開源(Open Source)許可證的風險與感染性
- (20 minutes)

University of Minnesota Banned from Linux Kernel Contributions (2021)

- **Research Experiment:** UMN researchers submitted patches with intentional vulnerabilities to the Linux kernel to study the feasibility of introducing stealthy bugs into open-source software. An IEEE S&P 2021 paper was accepted based on the research.
- **Community Reaction:** The Linux kernel maintainers were unaware of the experiment and felt their trust was violated, leading to significant backlash.
- **Consequences:** Greg Kroah-Hartman, a senior Linux kernel maintainer, banned UMN from contributing to the Linux kernel. All previous contributions from UMN were also removed and re-evaluated.
- **Apology Issued:** UMN researchers apologized, acknowledging the lack of prior consultation with the Linux community and the unintended harm caused. The authors retract the paper from IEEE S&P 2021.
- **Ethical Discussions:** The incident sparked broader discussions about ethics in open-source research and the responsibilities of contributors.


Recommend reading

- http://keithcu.com/wordpress/?page_id=599
- (Linux “chapter”, “After the Software Wars,” Keith Curtis, 2010)
- <http://keithcu.com/SoftwareWars.pdf>

Linux kernel in a diagram

Linux Kernel v2.6.11.8

"Woozy Beaver"




- 50% device driver
 - “Read the first 64 bytes of of /etc/passwd” —> “fetch block 3,348 from device 3”
- 25% CPU specific code
- Two inner layers are generic
- Mostly written in C + some assembly language to interface with hardware or chip specific functions

Advantages of Open Source Kernel


- Refactoring (smoothing, refining, simplifying, polishing) is done continuously in Linux
- Code is not freely available & in one place —> hard to evolve
 - Comparison: Microsoft's proprietary kernel —
The need to keep the old version alive for old code/hardware
 - Example: If many drivers have similar tasks, duplicate logic can be pulled out and put into a new subsystem that can then be used by all drivers.
- Stanford research: Linux kernel has **0.17 bugs per 1,000 lines of code, 150 times less** than average commercial code containing 20-30 bugs per 1,000 lines.

Ref: http://keithcu.com/wordpress/?page_id=599



<https://ma.ttias.be/system-calls-in-apache-linux-vs-iis-windows/>

System call graph in Microsoft's proprietary web server, **IIS**.



System call graph to return a picture in the
free web server **Apache**.

In-class: learn “screen” or “tmux”

- Allow you to have multiple “tabs” in terminal window
- Ctrl-a c to create a new “tab” (screen)
- Ctrl-a <number> to switch to that window (screen)
- Ctrl-a d to detach (screen)
- Re-attach using “screen -r”; useful when waiting for results
- tmux has different details but similar usage
- Use the workstation in the department or your VM to practice!

Drivers and Device Files

- Device driver: manages the interaction between the system and the hardware
- User space access: from /dev. Kernel maps operations “on these files” to calls to the driver code.
- Major and minor device numbers - (use ‘ls -l /dev’ to see them) map device file references to drivers
- Block device - read or write one block (multiples of 512) at a time
- Character device - read or write one byte at a time
- Some of the device driver’s functions
attach close dump ioctl open probe
prize read receive reset select stop
strategy timeout transmit write

“Phantom devices”

- `/dev/zero` & `/dev/null`:
data written here is discarded.
read from `/dev/zero` always returns 0.
read from `/dev/null` always returns EOF.
- `/dev/random` and `/dev/urandom`:
read from `/dev/random` will return random bytes.
(interface to kernel's random number generator)

Custom kernels v.s. loadable modules

- When installed, a system comes with a generic kernel
- Linux's udev system can manage real-time device changes
- Do we need custom-built kernels?
 - Pros: Opportunity for performance gain
 - Cons: Patch & system upgrade could be difficult
- For stability reasons:
using the stock kernel is recommended.

Kernel module related commands / files

- Files:
 - /vmlinuz: the actual kernel binary file
 - /lib/modules: kernel modules (separated by version)
- lsmod: list all kernel modules that have been loaded
- depmod: generate kernel module dependency file (modules.dep)
- modinfo: list information about a particular kernel module
- modprobe: load a kernel module (and its dependency)
- insmod, rmmod: manually load and unload kernel module

Initial ramdisk

- You might have seen this: `initrd.tbz`
- This is initial RAM disk (a disk in the memory)
- This is loaded before the root file system is mounted.
- Reason: some essential kernel modules might be needed for initial boot operation (storage, filesystem), but they are in `/lib/modules`
- Solution: load `initrd`, where these modules are stored.

Kernel Configuration

1. Modify tunable kernel configuration parameters

- Parameters can be adjusted via hooks in /proc (procfs)
- /proc/sys contains various special files for users to view and set kernel options
- Try it:
cat /proc/sys/fs/file-max
(maximum number of files the system can open at once)
sudo sh -c "echo <an integer> > /proc/sys/fs/file-max"
- Note that the changes are not carried across reboots

When to upgrade the kernel?

- Resist the temptation to keep up with the latest version
- Weigh between needs & risks
- Good rule of thumb:
Upgrade or apply patches only when expected productivity gains is larger than the effort to install.